

Akshay Raman

New York, NY
United States

phone: +1 (551) 804-2148
email: ar8692@nyu.edu
web: <https://akshayraman.com/>

Education

- **New York University, Courant Institute** New York, United States
Master of Science in Computer Science, AI Specialization Sept. 2023 - May 2025
 - GPA: 4.0/4.0
 - Capstone Project: Continual Credit Assignment in RL using Eligibility Traces. [Report](#)
- **Vellore Institute of Technology** Vellore, India
Bachelor of Technology in Computer Science and Engineering Sept. 2019 - July 2023
 - GPA: 9.28/10.00
 - Thesis: Neural Optimal Transport. [Report](#)

Technical Skills

Programming Languages: Python, C/C++, R, Java, SQL

Machine Learning Workflows: PyTorch, JAX, TensorFlow, Gymnasium, WandB, DeepSpeed, HuggingFace

Tools and Libraries: NumPy, SciPy, Pandas, OpenCV, OpenMP, CUDA, Docker, Git/GitHub, Linux

Research Experience

- **AI4VS Lab, Columbia University** with Prof. Kaveri Thakoor New York, United States
Machine Learning Researcher Jan. 2025 - Present
 - Investigating multimodal learning for clinical diagnosis using gaze patterns, retinal scans, and physician notes.
 - Exploring CLIP and vision-language models (VLMs) to combine visual, text, and gaze modalities for improved prediction.
- **DICE Lab, New York University** with Prof. Chinmay Hegde New York, United States
Graduate Research Assistant Sept. 2024 - Jan. 2025
 - Pursued research in data-centric ML advised by Prof. Chinmay Hegde at NYU Tandon School of Engineering.
 - Designed data curation pipelines to reduce spurious correlations in large-scale vision datasets, improving model generalization.
- **AI4Science Group, University of Ottawa** with Prof. Augusto Gerolin Ontario, Canada
Mitacs Globalink Research Intern June 2022 - Sept. 2022
 - Built deep neural solvers for high-dimensional Optimal Transport to simulate atom dissociation in Density Functional Theory (DFT). [Code](#)
 - Compared performance with traditional sinkhorn solvers, demonstrating faster convergence to approximate solutions.
 - Led interactive tutorials and seminars on ML foundations for non-CS students.

Publications

- [1] Khan, R. et al. (including Raman, A.) “Use of artificial intelligence algorithms to predict systemic diseases from retinal images” *WIREs Data Mining and Knowledge Discovery*, Vol. 13, No. 5 (2023)

Projects

1. Hierarchical CLIP-based Image Geolocation Prediction [Link](#)

- Trained a custom CLIP-based model on 4M+ MediaEval images, achieving 70% country-level accuracy in image geolocation.
- Reduced inference latency and memory usage by $\sim 100\times$ via hierarchical feature clustering, enabling real-time predictions.

2. Consistent Multi-view Object Generation using Finetuned Video Diffusion Models [Link](#)

- Fine-tuned a video diffusion model to generate multi-view consistent object renderings from single-view inputs.
- Demonstrated that a curated high-quality 1% subset (10K objects) of the Objaverse dataset achieved performance comparable to full-scale training. (1M+ objects).

3. Diabetic Retinopathy Detection [Link](#)

- Trained large-scale CNNs to predict diabetic retinopathy (an eye disease) from a noisy dataset of retinal images.
- Used Grad-CAM to generate saliency maps and visualize regions influencing model predictions.

Teaching Experience

NYU CSCI-UA.0202 Operating Systems <i>Course Assistant</i>	Spring 2025
NYU CSCI-GA.3033 Graphical Processing Units (GPUs) <i>Grading Assistant</i>	Fall 2024
NYU CSCI-GA.3033 Multicore Processors <i>Grading Assistant</i>	Spring 2024
NYU CSCI-UA.0480 Parallel Computing <i>Grading Assistant</i>	Fall 2023